

# OS Lab 1: System Calls and Schedulers

**Part 1:** Add a system call `info(int param)` to return number of processes running and number of system calls made.

**List of files modified and functions added:**

1. `usys.S` - Add line `SYSCALL(info)`
2. `user.h` - Add line `int info(int);`
3. `syscall.h` - Add line `#define SYS_info 22`
4. `syscall.c` - Add lines `extern int sys_info(void)` and `[SYS_info] sys_info` and in `syscall` method `proc->call_count++;`
5. `sysproc.c` - Add snippet

```
int sys_info(void){
    int n;
    if( argint(0, &n) < 0)
        return -1;
    return info(n);
}
```
6. `proc.c` - Add the system call logic

```
int info(int param){
    int count = 0;
    struct proc *p;
    switch(param){
    case 1:
        for(p = ptable.proc; p<&ptable.proc[NPROC]; p++){
```

```

        if(!(p->state == UNUSED || p->state == ZOMBIE))
            count += 1;
    }
    printf("No of running processes: %d\n", count);
    break;
case 2:
    printf("No of syscalls made by the process: %d\n", proc->call_count);
    break;
case 3:
    printf("No of memory pages used: %d\n", proc->page_count);
    break;
}
return count;
}

```

In allocproc method add

```

p-> pages_count = 0;
before kalloc is called
p->page_count++;
p->call_count = 0;

```

In fork method, when kfree is called

```

add if(np->page_count){
    np-> page_count - -;
}

```

In wait method add

```

if(p->page_count){
    p-> page_count - -;
}

```

7. proc.h - Add line in struct proc

```

    int call_count;
    int page_count;

```

8. def.h - Add line `int info(int);`
9. test.c - Test function to check the system call

```
#include "types.h"
#include "stat.h"
#include "user.h"

int main(int argc, char *argv[]) {
    info(1);
    info(2);
    info(3);
    exit();
}
```

10. Makefile - Add line `_test\`

## Results :

```
eter -fno-stack-protector -c -o proc.o proc.c
ld -m elf_i386 -T kernel.ld -o kernel entry.o bio.o console.o exec.o file.o fs.o ide.o idt.o kalloc.o
kbd.o lapic.o log.o main.o mp.o picirq.o pipe.o proc.o sleeplock.o spinlock.o string.o switch.o syscall.o s
ysfile.o sysproc.o timer.o trapasm.o trap.o uart.o vectors.o vm.o -b binary initcode entry.o
objdump -S kernel > kernel.asm
objdump -t kernel | sed '1,/SYMBOL TABLE/d; s/ .* / /; /^$/d' > kernel.sym
dd if=/dev/zero of=xv6.img count=10000
10000+0 records in
10000+0 records out
512000 bytes (5.1 MB) copied, 0.0311333 s, 164 MB/s
dd if=bootblock of=xv6.img conv=notrunc
1+0 records in
1+0 records out
512 bytes (512 B) copied, 0.000236396 s, 2.2 MB/s
dd if=kernel of=xv6.img seek=1 conv=notrunc
135+1 records in
135+1 records out
171847 bytes (172 kB) copied, 0.03114081 s, 151 MB/s
qemu -nographic -drive file=fs.img,index=1,media=disk,format=raw -drive file=xv6.img,index=0,media=disk,for
mat=raw -smp 2 -m 512
Could not open option rom 'sgabios.bin': No such file or directory
xv6...
cpu0: starting
cpu0: starting
sb: size 1024 nblocks 941 ninodes 288 nlog 32 logstart 2 inodestart 32 heap start 58
init: starting sh
$ test
No of running processes: 3
No of syscalls made by the process: 4
No of memory pages used: 1
$
```

## PART 2 :

### 1. Lottery Scheduling :

List of files that were modified :

- usys.h
- user.h
- syscall.h
- syscall.c
- sysproc.c
- proc.c
- defs.h

Snippets of the changes made in proc.c are as follows :

```
//our code for set_tickets
int
set_tickets(int tickets)
{
    total_tickets += tickets;
    proc->tickets = tickets;
    cprintf("ticket value set to :%d\n", proc->tickets);
    return 0;
}
```

System call to set the value of tickets

The lottery scheduler code :

```
void
scheduler(void)
{
    struct proc *p;
    unsigned long winner;
    int counter;
    int sd=0;

    for(;;){
        // Enable interrupts on this processor.
        sti();
        sd++;
        //OS 202
        //Random number generator
        //Using lottery

        winner=rdm_gen(sd)*(total_tickets+1);

        //cprintf("Winner is : %d\n", winner);
        counter = 0;
        // loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            //cprintf("Insides %d", p->state);
            if(p->state != RUNNABLE)
                continue;
            //cprintf('Counter: %d, Tickets: %d', counter, p->tickets);
            counter = counter + p->tickets;
            if(counter < winner)
                continue;
        }
    }
}
```

```

        // Switch to chosen process. It is the process's job
        // to release ptable.lock and then reacquire it
        // before jumping back to us.
        //printf("pid : ", p->pid);
        p->count += 1;
        pcpu = p;
        switchuvm(p);
        p->state = RUNNING;
        switch(&cpu->scheduler, p->context);
        switchkvm();

        // Process is done running for now.
        // It should have changed its p->state before coming back.
        proc = 0;
    }
    release(&ptable.lock);
}
}

```

The random number generator :

```

/*
unsigned long
rndm_gen(int sd)
{
    return (sd*279470279UL)%4294967291UL;
}

```

Changes made in allocproc() function :

```

// Set up new context to start executing at forkret,
// which returns to trapret.
sp -= 4;
*(uint*)sp = (uint)trapret;

sp -= sizeof *p->context;
p->context = (struct context*)sp;
memset(p->context, 0, sizeof *p->context);
p->context->esp = (uint)forkret;
p->tickets = 1;
p->count = 0;
p->pass = 1;
//total_tickets -= 1;
return p;
}

```

The variable tickets is defined in struct proc and given a default value in allocproc(). Also, the variable count, which calculates the number of counts a process is scheduled to run by the scheduler is given a default value of 0.

Changes made in sysproc.c :

```

//Our system call
int
sys_set_tickets(void)
{
    int n;
    if(argint(0, &n) < 0)
        return -1;
    //if(argint(1, &pid) < 0)
    //    return -1;

    return set_tickets(n);
}

```

Results obtained on running processes with 3 tickets :

```
[ $ test8; test28; test1
number of runs of sh is 2
number of runs of sh is 2
ticket value set to :20
number of runs of test1 is 8
$ ticket value set to :30
number of runs of test is 6
zombie!
ticket value set to :10
number of runs of test2 is 6
zombie!
```

## 2. Stride Scheduler :

List of files that were modified :

- usys.h
- user.h
- syscall.h
- syscall.c
- sysproc.c
- proc.c
- defs.h

Snippets of the changes made in proc.c are as follows :

```
void
scheduler(void)
{
    struct proc *p;
    int min_pass;
    int w_pid=0;

    for(;;){
        min_pass = 10000;
        // Enable interrupts on this processor.
        sti();
        //OS 262
        // loop over process table looking for process to run.
        acquire(&table.lock);

        for(p = table.proc; p < &table.proc[NPROC]; p++){
            //cprintf("insides: %d", p->state);
            if(p->state != RUNNABLE)
                continue;
            if(p->pass < min_pass){
                min_pass = p->pass;
                w_pid = p->pid;
            }
        }
        for(p = table.proc; p < &table.proc[NPROC]; p++){
            //cprintf("insides: %d", p->state);
            if(p->state != RUNNABLE)
                continue;
            if(p->pid != w_pid)
                continue;
        }
    }
}
```

```

// Switch to chosen process. It is the process's job
// to release ptbls.lock and then reacquire it
// before jumping back to us.
//printf("pid : ",p->pid);
p->ccurt +=1;
p->pass -= (1000/p->tickets);
proc = p;
switchproc(p);
p->state = RUNNING;
switch(&proc->scheduler, p->context);
switchken();

// Process is done running for now.
// It should have changed its p->state before coming back.
proc = 0;
}
release(&ptbls.lock);
}
}

```

Results obtained :

```

[$ test&;test1&;test2&
number of runs of sh is 23
number of runs of sh is 2
ticket value set to :30
number of runs of test is 5
zombie!
ticket value set to :20
number of runs of test1 is 4
number of runs of sh is 9
zombie!
$ ticket value set to :10
number of runs of test2 is 11
zombie!

number of runs of sh is 1

```